

 **Question 1 Correct**

^

You want to move your Hadoop on-prem implementation to the cloud, to support future increase in traffic.

Which cloud service should you use for that?

EventArc

GKE

Your answer is correct

 **Dataproc**

Overall explanation

#1 is incorrect. GKE does not provide native Hadoop platform, and running it on GKE cluster will not necessarily provide the required scale

#2 is incorrect since EventArc is not Hadoop compatible

#3 is correct. Dataproc allows running Hadoop in the cloud in a fully managed, easy to use and low cost service

 **Question 2 Correct**

You have an App Engine running in the `europe-west2` region. You want to deploy an app to the App Engine that will run in the `us-east1` region.

What should you do?

Your answer is correct

- Create a new Project, create a new App Engine in the `us-east1` region and deploy the new app there**

- Create a new version of the app and configure it to run in the `us-east1` region**

- Create a new service in the existing App Engine and configure it to run in the `us-east1` region**

Overall explanation

#1 is correct. You can't have App Engine apps in more than one region in a single project

#2 is incorrect since you cannot configure a service to run in a different region

#3 is incorrect since you cannot configure a version to run in a different region

 **Question 3 Correct**

^

You need to design a cloud-based system that will ingest petabytes of data for analytics. The data should be accessible using SQL interface.

Which datastore solution should you use?

Cloud SQL

Your answer is correct

 **BigQuery**

Cloud Storage

Overall explanation

#1 is incorrect since Cloud Storage does not have SQL interface

#2 is incorrect since Cloud SQL cannot handle petabytes of data

#3 is correct since BigQuery is designed for exactly these scenarios

Question 4 Incorrect

^

True or false:

It's a good practice to assign users who should not be able to modify resources the `roles/viewer` role.

Correct answer

False

Your answer is incorrect

True

Overall explanation

This statement is False.

You should not use the basic roles (`roles/viewer`, `roles/editor`, `roles/owner`) with users, and they are too broad. Assign only resource-specific roles when possible.

Question 5 Correct

You have an app running on an App Engine service named "ecommerce". Currently there's a single instance running in this service.

Black Friday is getting close, and you want to make sure the app will be able to scale as needed and add instances to support the expected load.

What should you do?

Your answer is correct

- Configure scaling using the `app.yaml` configuration file

- Set up managed instance group and configure scaling

- Migrate the app to Cloud Run which supports auto scaling

Overall explanation

#1 is incorrect since App Engine supports scaling

#2 is correct. The `app.yaml` file is the place to configure scaling for App Engine

#3 is incorrect. Managed instance group is a feature of Compute Engine and not of App Engine

 **Question 6 Correct**

You're designing an app for storing sensitive HR data. The data is stored in a Cloud Storage.

You want to make sure the information will not be visible to users who should not see it.

What do you do in order to limit access to the data?

Set up Private Google Access to the Cloud Storage

Encrypt the data in code before uploading it to the Cloud Storage bucket

Your answer is correct

 **Assign the `storage.objectViewer` role only to users and service accounts that need to use the data.**

Overall explanation

#1 is correct. By granting `storage.objectViewer` role only to users and service accounts that require access to the data, you deny access from parties who do not need it.

#2 is incorrect. Encrypting the data does not limit its access.

#3 is incorrect. Setting up Private Google Access will not prevent users from accessing the data through the console or CLI.

 **Question 7 Correct**

Your company wants to archive 50TB of its data. The data will not be accessed frequently, but should be stored for a long term.

The company decided to check whether the cloud can be a solution for this archive, while minimizing costs.

How would you implement this archive?

- Using Log Analytics**

Your answer is correct

-  **Using the Archive tier of Cloud Storage**

- Using the lowest tier of Cloud SQL**

Overall explanation

#1 is incorrect since Cloud SQL will be quite expensive for this amount of data

#2 is incorrect since Log Analytics cannot be used to store external data

#3 is correct. The Archive tier of Cloud Storage targets this scenario exactly

 **Question 8 Correct**

^

Your company has a system that gets a stream of data generated from sensors deployed across the country. You want to analyze and enrich this data, while minimizing maintenance required from cloud engineers.

Which cloud service should you use?

Your answer is correct

 **Dataflow**

BigQuery

Pub/Sub

Overall explanation

#1 is incorrect since Pub/Sub does not analyze data, just publish messages

#2 is incorrect since BigQuery does not work with stream of data

#3 is correct since Dataflow is designed exactly for these scenarios



Question 9 Correct

You set up a global external load balancer in front of a managed instance group serving a web app. When trying to access the web app through the load balancer you get an error, and when looking at the firewall page in the console you see the instance group is marked as "Unhealthy".

What can be done to solve this problem?

- Check the code deployed in the instance group for any errors
- Make sure the instance group is configured to allow access to HTTPS traffic only

Your answer is correct

- Add firewall rule to the instance group, allowing access from 35.191.0.0/16 and 130.211.0.0/22 IP ranges

Overall explanation

#1 is incorrect since code errors will not make the instances unhealthy

#2 is incorrect since the health of the instances are not related to the protocol served by them

#3 is correct since these are the addresses used by the health probes of the load balancer, and if the probes cannot access the instances they will be marked as "Unhealthy"



Question 10 Correct

You design a global system in the cloud. The system should be as performant as possible, without incurring excessive costs.

You want to find a solution for serving static content to users as fast as possible. Which cloud service should you use?

Your answer is correct

Cache the content in Cloud CDN

Multiple deployments of Cloud Run, each deployed in a different region

Cloud Storage buckets in multi-region configuration

Overall explanation

#1 is incorrect since not all buckets will provide the same latency

#2 is incorrect since the solution is going to be quite expensive

#3 is correct. Cloud CDN targets this exact scenario, and should be used to bring static content closer to the end user



Question 11 Correct

Your web app has two VMs, one for front end and one for batch processing. You need to make sure access to the batch VM is allowed only from the front end VM, and that the front end VM is open to the internet.

How do you configure it?

- Configure Cloud Armor on the VPC of the front end VM and deploy the batch VM as a managed instance group
- Place an external load balancer in front of the front end VM and an internal load balancer in front of the batch VM

Your answer is correct

- Configure firewall rules to allow external traffic to the front end VM and to block traffic to the batch VM that was not originated from the front end VM

Overall explanation

#1 is incorrect since load balancers are used to balance load between instances, and not to deny or allow network routing

#2 is correct. This is the best way to achieve that

#3 is incorrect. Cloud Armor is used to protect against web attacks and not to route traffic, and managed instance group will help with scaling, but not in restricting traffic



Question 12 Correct



True or false:

You should never use more than one VPC in a single project

True

Your answer is correct

False

Overall explanation

This statement is False.

It's usually a good idea to use more than one VPC in a project in order to better segment the infrastructure and use a hub-and-spoke model.

 **Question 13 Correct**

^

You design a gaming system in the cloud that should serve million of users world wide.

The system is based on GKE cluster.

In order to make the system as reliable as possible, you decide to deploy two clusters, one in the US and one in Europe. The clusters run the same codebase.

What should be added to the system in order to automatically route users to the closest GKE cluster, and to detect when a cluster goes down and then route users to the healthy one?

Your answer is correct

Global Load Balancer

Spanner in Multi-Region configuration

Cloud CDN

Overall explanation

#1 is correct. Global Load Balancer should be used for these exact scenarios, as it can route requests to healthy backends across regions.

#2 is incorrect, as CDN does not route requests to backends, but caches static content

#3 is incorrect. While Spanner can be configured to support multi-region scenarios, it's a database and not a backend running code



! Question 14 Incorrect



You want to deploy a new version of an existing app in App Engine. You want testers to validate the new version before releasing it to the public, and you want to make sure current version will still use the same resources and URL so that user experience will not be hurt.

What is the type of deployment you're looking for?

Rolling

Your answer is incorrect

Canary

Correct answer

Blue-Green

Overall explanation

#1 is incorrect since the same resources are used for the existing and new version

#2 is correct since new resources are allocated to the new version

#3 is incorrect since the same resources are used for the existing and new version

 **Question 15 Correct**

^

You design an agriculture IoT system that processes telemetry received from sensors deployed in fields around the world. The internet connection from these sensors is unstable.

You need to design a solution that will maximize the reliability of the data ingestion while minimize costs.

- Each sensor will establish a connection to cloud function, and will stream the messages to it**

Your answer is correct

- The sensors will send the messages to Pub/Sub, which then will publish them to the respective listeners**

- Each sensor will have its own cloud storage bucket, and will send the messages to it**

Overall explanation

#1 is correct. Pub/Sub is a very efficient mechanism to publish messages to other services, and no persistent connection is required

#2 is incorrect since we cannot trust the connection to remain stable

#3 is incorrect since it does not include notifying other services on the new data, and also has higher costs

 **Question 16 Correct**

^

Which of the following CANNOT be used as a destination for EventArc?

Your answer is correct

 **App Engine**

GKE

Cloud Functions

Overall explanation

The correct answer is #2.

App Engine cannot be used as a destination for EventArc, while Cloud Functions and GKE can.



Question 17 Correct



You want to deploy your .NET based app in GCP.

Which of the following compute resources can't be used for that?

Cloud Run

Your answer is correct

App Engine Standard

Compute Engine

Overall explanation

#1 is incorrect. You can deploy whatever you want on a virtual machine

#2 is correct. App Engine Standard does not support .NET

#3 is incorrect. You can deploy .NET apps in Cloud Run



Question 18 Correct

^

Your client deployed a new version of an app in App Engine. After a few hours he started to receive reports of slowness in the new version.

What should be the first step to take in this scenario?

- Migrate to Cloud Run

Your answer is correct

- Roll back to the previous version

- Move to a stronger instance class

Overall explanation

#1 is correct. Roll back to the previous version which did not have this problem, and then research the problem

#2 is incorrect. We can't know at this stage that the problem is due to a lack of resources

#3 is incorrect. We can't be sure this will solve the problem, and such migration takes time

Question 19 Incorrect

The development team in the project you're working on wants to move development to cloud-based VMs using Compute Engine. You want to make sure cost is minimized.

What should the team do in order to minimize cost?

Your selection is correct

- Use Instance Schedules to set stop and start times for the virtual machines**

Your selection is incorrect

- Use Managed Instance Group to scale the machines as needed**

Correct selection

- Make sure to place the code on a persistent disk**

Your selection is incorrect

- Make sure to place the code on a local SSD**

Overall explanation

#1 is correct since by placing the code on a persistent disk the code will be preserved across machine stop and start

#2 is incorrect since the code will be lost when the machine stops

#3 is incorrect since we don't need to scale development machines

#4 is correct since Instance Schedules save costs by turning off the machines when not needed



Question 20 Correct



You design a system in the cloud that should support multi-region writes in its database.

Which of the following databases provides such functionality?

Your answer is correct

Spanner

Cloud SQL

AlloyDB

Overall explanation

#1 is incorrect. Cloud SQL does not provide multi-region writes

#2 is correct. Spanner provides multi-region writes

#3 is incorrect. AlloyDB does not provide multi-region writes

 **Question 21 Correct**

^

You plan to deploy a new app in the cloud. One of the most important features you're looking at is traffic splitting, which will allow you to test new versions without interrupting existing versions.

Which of the following compute platform support this feature?

Compute Engine

Your selection is correct

Cloud Run

Your selection is correct

App Engine

GKE

Overall explanation

#1 is incorrect. Compute Engine provides general purpose virtual machine, and does not have the notion of "traffic"

#2 is correct. App Engine supports traffic splitting using versions

#3 is correct. Cloud Run supports traffic splitting using revisions

#4 is incorrect. GKE does not support traffic splitting

 **Question 22 Correct**

^

You want to migrate a legacy application from an on-premises data center to GCP. The application currently runs on virtual machines. You're not going to make any changes to the code, but you want to make sure the application can handle expected load, and that it's protected against malicious attacks.

What would you do?

- 1. Deploy the app on multiple Compute Engine instances in the cloud**
- 2. Place the instances behind an external load balancer**
- 3. Configure VPC rules to protect the incoming traffic**

Your answer is correct

- 1. Deploy the app on multiple Compute Engine instances in the cloud**
- 2. Place the instances behind an external load balancer**
- 3. Configure Cloud Armor on the load balancer**

- 1. Deploy the app on Google Kubernetes Engine cluster in the cloud**
- 2. Configure firewall rules for the GKE**

Overall explanation

#1 is incorrect since VPC rules do not protect against malicious internet attacks, just against disallowed traffic between cloud resources.

#2 is correct, since Cloud Armor is the mechanism to use to protect against malicious web attacks.

#3 is incorrect since we don't want to change the code, and that means we'll probably won't be able to deploy it on GKE as is

! Question 23 Incorrect

^

Your client wants to connect its on-prem network to a VPC in Google Cloud. The connection should be secure and fast.

Which two steps should the client take in order to successfully connect the two networks?

Correct selection

- Create a new VPC with IP range that's not overlapping the one used by the on-prem network

Your selection is incorrect

- Set up a VPN connection between the on-prem network and the new VPC

- Create a new VPC with Auto Mode

Correct selection

- Set up a Cloud Interconnect between the on-prem network and the new VPC

Overall explanation

#1 is correct. You cannot peer two networks with overlapping IP range

#2 is incorrect. When creating a VPC in auto mode you don't have control on the IP range and it might overlap the IP range of the on-prem network

#3 is incorrect. VPN is not the most secure and fast connection between on-prem and the cloud

#4 is correct. Cloud Interconnect is the most secure and fast connection between on-prem and the cloud

 **Question 24 Correct**

^

You developed a small app for managing inventory, and want to deploy it in the cloud. Your budget is tight, and you want to use a platform that won't charge anything when the app is not in use.

Which platform should you use?

GKE

App Engine Flexible Environment

Your answer is correct

 **App Engine Standard Environment**

Overall explanation

#1 is correct since you can scale to 0 with the Standard environment

#2 is incorrect since you'll pay even if the app is not used

#3 is incorrect since you'll pay even if the app is not used

 **Question 25 Incorrect**

^

Your client has an API that handles data request from a Cloud SQL. The API is implemented using App Engine, and is protected using an external load balancer.

The client is going to release a new version of the API, and he wants to allow a subset of developers to test the new version. These developers should use the same URL they're using today in order to access the new API version.

How would you implement this requirement?

Correct answer

- Add a new backend to the load balancer with the new version**

Your answer is incorrect

- Deploy a new version of the App Engine**

- Deploy a new load balancer pointing to the new App Engine**

Overall explanation

#1 is incorrect, since versions in App Engine have their own URL, and the requirements clearly state the developers should use the same URL

#2 is correct. The load balancer can be configured to route based on the source IP address

#3 is incorrect. A new load balancer will have its own URL

 **Question 26 Correct**

^

You design a data intensive IoT application in the cloud, and need to decide which database to use in the system.

Which of the following correctly represents a reason to select Spanner as the database?

- Spanner costs less than CloudSQL**

Your answer is correct

- Spanner supports horizontal scaling and can deal with large amount of data**

- Spanner supports SQL and NoSQL data formats**

Overall explanation

#1 is incorrect, as Spanner usually costs more than CloudSQL.

#2 is incorrect. Spanner is a relational database

#3 is correct. Spanner supports horizontal scaling

 **Question 27 Correct**

^

You review a data intensive cloud system. One of the components of the system, running in Cloud Run, works with Cloud SQL to store transactional data. Looking at the component's code you find out that the code contains the connection string to the Cloud SQL.

What would be your recommendation regarding that?

- Store the connection string in an external configuration file**
- Keep it as it is. The code is secure within the Cloud Run runtime, and there's no problem storing there connection strings**

Your answer is correct

- Store the connection string in Secrets Management**

Overall explanation

#1 is incorrect. You should never store secrets in code as it's less secure than the other options.

#2 is incorrect since other can access the configuration file and read the connection string

#3 is correct. Secrets Management is the best place to store such secrets, and offers the best security for the data

 **Question 28 Correct**

^

You develop a system that sends stock quotes to subscribers. You want to use a cloud service that will be used to send these messages and allow subscribers to retrieve them as quickly as possible.

You don't know how many subscribers there's going to be.

Which service should you use?

Your answer is correct

 **Pub/Sub**

Dataflow

Dataproc

Overall explanation

#1 is incorrect. Dataflow should not be used to publish messages to subscribers

#2 is incorrect. Dataproc is not used for this scenario

#3 is correct. Pub/Sub offers the Fan-out pattern which can be used to send messages to many subscribers

 **Question 29 Correct**

When deploying your solution to production, you repeatedly find bugs in the production that were not found on the test environment.

What should you do?

- Add monitoring to the code**

Your answer is correct

-  **Deploy changes to small subset of users to test before the roll-out**

- Conduct a security review on the code**

Overall explanation

#1 is correct. It's best to have the system tested in its final environment before rolling it out to all the users

#2 is incorrect. The bugs are not security related

#3 is incorrect. While it's always a good idea to have monitoring in the system, you need it to be tested before roll out

 **Question 30 Correct**

You have a Compute Engine instance in a VPC named `vpc1`. There's another instance, placed in a VPC named `vpc2`.

As part of a new feature development of the system, these two instances should be able to communicate with each other, and you decide to use peering for that.

What you should ensure before setting up the peering?

Your answer is correct

That the two VPCs do not have overlapping IP range

That at least one VPC is protected using an internal firewall

That the two VPCs are in the same region

Overall explanation

#1 is incorrect since VPCs are global resources, and do not reside in a specific region

#2 is correct. You cannot set up peering between VPCs with overlapping IP range

#3 is incorrect. You don't have to have firewall in order to set up peering

 **Question 31 Correct**

^

You deploy a Cloud Function which responds to incoming HTTP requests. You want to make sure the function response time is as quick as possible, and one of the things you want to avoid is cold start.

How can you avoid that?

Your answer is correct

- Set minimum number of instances**

- Set maximum number of instances**

- Add memory to the function**

Overall explanation

#1 is incorrect. Adding memory will not prevent cold start

#2 is incorrect. Adding instances will not prevent cold start after all instances have shut down

#3 is correct. Setting minimum number of instances ensures there's always a certain number of running instances

 **Question 32 Correct**



True or false: When deploying Cloud Run from source, no Artifact Registry is used

Your answer is correct

 **False**

True

Overall explanation

This statement is False.

When deploying from source, an Artifact Registry named `cloud-run-source-deploy` is created and stores the code to deploy

 **Question 33 Correct**

^

Your client wants to migrate a system running on C++ to the cloud. The system is very old, and no one in the company knows how to handle it, which means no code changes can be made.

What is the compute platform recommended for this migration?

Cloud Run

Your answer is correct

 **Compute Engine**

App Engine

Overall explanation

#1 is incorrect since App Engine does not support C++

#2 is incorrect since you can't always containerize C++ code

#3 is correct since this option gives you the most control and flexibility when deploying the code

 **Question 34 Incorrect**

^

Your client has an on-prem datacenter with Oracle database. The client wants its App Engine-hosted app to be able to connect to this database, without exposing the database to the internet.

The on-prem datacenter is quite far away from any Google Cloud datacenter.

What would you recommend the client to do?

- Set up peering between the App Engine app and the on-prem datacenter

Correct answer

- Set up Partner Interconnect connection between the datacenter and Google Cloud

Your answer is incorrect

- Set up Dedicated Interconnect connection between the datacenter and Google Cloud

Overall explanation

#1 is incorrect since it's not possible to set up peering between on-prem network and the cloud without any other connection.

#2 is incorrect since the on-prem datacenter is far away from any Google Cloud datacenter

#3 is correct, and provides the best method for connecting the on-prem datacenter to Google Cloud.

 **Question 35 Correct**

Your Java app is running on a Compute Engine instance in the cloud. As part of a new feature request, it was decided that a new BigQuery database will be used to provide analytics on metrics generated by the app.

You add the Java BigQuery client library to your code.

When running the code in the cloud, you get an error stating you don't have permissions to connect to BigQuery.

What should you do?

- Configure BigQuery to allow connections from the internet**

Your answer is correct

- Use a Service Account in the code that has permissions to connect to BigQuery**

- Move the code to GKE**

Overall explanation

#1 is correct since using Service Account is the best practice to configure authentication to Google services

#2 is incorrect since moving to another compute platform will not solve the authentication issue

#3 is incorrect since it's extremely unsafe to expose a database to the internet

 **Question 36 Correct**

^

You have an IoT system in the cloud that receives telemetry from weather sensors with data about temperature, humidity, location and more.

You need a cloud solution for storing and querying this data. The solution should minimize costs.

What should you do?

Store the telemetry in Cloud Storage

Store the telemetry in Log Analytics

Your answer is correct

 **Store the telemetry in BigQuery**

Overall explanation

#1 is incorrect since Cloud Storage, while cost effective, does not have querying capabilities.

#2 is correct. BigQuery has a cost effective storage costs and great querying capabilities.

#3 is incorrect since Log Analytics cannot be used to store external data.

! Question 37 Incorrect

^

You have a Compute Engine instance running a publicly accessible web app. The instance has a locally-installed database, and as the system grows you decide it's time to move to a fully managed database and you migrate the data to Cloud SQL.

You now need to make sure the instance connects to the Cloud SQL using private IP, and not using public IP which is exposed to the internet.

Which mechanism should you use for that?

Your answer is incorrect

Private Google Access

Correct answer

Private Service Access

Private Service Connect

Overall explanation

#1 is incorrect, since Cloud SQL is not supported for Private Google Access

#2 is incorrect, since Cloud SQL is not supported for Private Service Connect

#3 is correct. Cloud SQL can be configured to be connected to using Private Service Access



Question 38 Correct



You design a system that receives usage details of devices from around the world. The data should be stored for future analytics.

In average, the data rate will be 5000 messages / minutes, with bursts of up to 7000 messages / minutes.

Which storage solution should you use?

Your answer is correct

BigTable

Cloud Storage

Cloud SQL

Overall explanation

#1 is correct. BigTable provides high-throughput analytics and native time series support that are usually needed for data analytics, and it can stand in the data rate requested.

#2 is incorrect, since data stored in Cloud Storage cannot be queried

#3 is incorrect since Cloud SQL is not designed to withstand such a rate



Question 39 Correct

You have a Compute Engine instance named `hr-vm1` running in a VPC named `hr-vpc`. The instance is running in the `us-east1` region, and does not have public IP address.

You also have App Engine deployed in the `europe-west3` region, with a service named `hr-app`.

The `hr-app` needs to communicate with the `hr-vm1`.

The communication should be as secure as possible and minimize costs.

What are the steps that should be taken for that?

- Set up a public IP in the `hr-vm1`, and call it from the `hr-app` service

Your answer is correct

- Set up Serverless VPC Access from the App Engine to the VPC and call the private IP of `hr-vm1`

- Set up a peering between the App Engine's VPC and the instance's VPC

Overall explanation

#1 is incorrect. App Engine do not reside in a VPC and therefore we cannot peer it to other VPCs

#2 is incorrect. We want the solution to be secure, and setting up public IP will compromise the security

#3 is correct. Serverless VPC Access provides secure access from serverless resources, such as App Engine, to VPC resources

! Question 40 Incorrect

You design a new containers-based gaming platform in GCP, and you want to make sure the system follows Google's security best practices.

What should you do?

Your selection is correct

- Ensure there's monitoring in place on the containers

Correct selection

- Ensure you're using native logging

Your selection is correct

- Avoid using privileged containers

Your selection is incorrect

- Ensure there are alerts in place when GKE monthly cost goes above a specific threshold

Overall explanation

#1 is correct since if attacked, privileged containers allow attacker to do a lot of damage in the container

#2 is correct since native logging uses the native cloud logging mechanism which is more streamlined and secure

#3 is incorrect since, while monitoring costs is important, it's not a security-related aspect

#4 is correct since with monitoring you can locate security breaches on your system

 **Question 41 Correct**

^

Which of the following is an optimal SLO definition?

- The website will have at least 99% availability**
- If the system will not meet the performance target, the supplier will pay a \$10K compensation to the client**

Your answer is correct

-  **No more than 2% of requests return 500 error code in a month**

Overall explanation

1# is correct, since it's measurable and concrete measure.

2# is incorrect since it's not clear what is "availability"

3# is incorrect since this is not a service level objective, but a service level agreement (SLA)

 **Question 42 Incorrect**

^

You have a Compute Engine instance running a publicly accessible web app. The instance has a locally-installed database, and as the system grows you decide it's time to move to a fully managed database and you migrate the data to Spanner.

You now need to make sure the instance connects to the Spanner using private IP, and not using public IP which is exposed to the internet.

Which mechanism should you use for that?

Correct answer

Private Service Connect

Private Service Access

Your answer is incorrect

Private Google Access

Overall explanation

#1 is incorrect, since Private Google Access cannot be used with instances which have public IP address

#2 is correct. Private Service Connect allow private connection between publicly accessible instances and spanner

#3 is incorrect. Private Service Access does not support Spanner

 **Question 43 Incorrect**

You have a Compute Engine instance that's becoming too small for your needs. Which of the following can be done on the instance without any downtime?

Correct answer

- Resizing disk**

Your answer is incorrect

- Adding memory**

- Changing instance class**

Overall explanation

#1 is incorrect since adding memory requires a restart

#2 is incorrect since changing instance class requires a restart

#3 is correct. Resizing instance disk does not require a restart



Question 44 Correct

Your Java app, running on Cloud Run, uses quite a few packages, and you want to store them in a central location in the cloud that has integration with your app.

What should you use for that?

Your answer is correct

Artifact Registry

Container Registry

GitHub

Overall explanation

#1 is incorrect. While you can store packages on GitHub, it's not integrated with Cloud Run

#2 is correct. Artifact Registry can be used to store Java packages, and it integrates with Cloud Run

#3 is incorrect. Container Registries are used to store container images, and not software packages

! Question 45 Incorrect

Your organization has multiple projects in GCP, all under the same organization hierarchy.

You want resources from Project1 to be able to connect to resources in Project2, with minimal cost and maintenance.

What should you do?

Your answer is incorrect

- Deploy VPC in every project and set up peering between the VPCs in the projects**

Correct answer

- Set up Shared VPC and define the projects as Service Projects**

- Deploy subnet in every project and set up peering between the subnets**

Overall explanation

#1 is correct. Shared VPC allows multiple projects in an organization to use the same VPC

#2 is incorrect. Setting up and peering VPCs require set up work, and we want to avoid it

#3 is incorrect. You cannot set up peering between subnets



Question 46 Correct

As part of an external audit conducted on your cloud system, the auditors need to have full visibility on the project's resources.

What is the role the auditors should be assigned in order to do that?

Project Owner

Your answer is correct

Project Viewer

Project Browser

Overall explanation

#1 is correct. The Project Viewer role grants access to view the resources in the project without changing them

#2 is incorrect. The Project Browser role allows access to the project hierarchy, including the folder and organization, but does not allow viewing resources

#3 is incorrect. The Project Owner allows too broad access, including creating and modifying resources



! Question 47 Incorrect



You have a system for processing telemetry from external sensors. The telemetry is sent using the TCP protocol, and received by a Compute Engine instance deployed in the `europe-west3` region.

You want to protect the incoming traffic and use a firewall. Which firewall type should you use?

- Application, global, external

Correct answer

- Proxy, regional, external

Your answer is incorrect

- Proxy, internal

Overall explanation

#1 is correct. The Proxy load balancer listens to TCP traffic, traffic is external, and the app is regional

#2 is incorrect. Application load balancer works with HTTP, and this scenario requires TCP

#3 is incorrect. The traffic is external and not internal



Question 48 Correct

You design an ecommerce website for a client. The performance and reliability of the website are extremely important to the client, and he asks you to define thresholds for them.

What you're asked to define is:

SLA

Your answer is correct

SLO

SLI

Overall explanation

#1 is incorrect. SLA, or Service Level Agreement, defines what happens if you fail to achieve the thresholds defined.

#2 is correct. SLO, or Service Level Objectives, is the target you aim to achieve with the system. For example - average request response time of 100ms.

#3 is incorrect. SLI, or Service Level Indicator, is the actual performance of the system, and how it's measured against the SLO.

! **Question 49 Incorrect**

^

You want to migrate existing code base to the cloud.

Which compute platform will make the migration the easiest?

GKE

Correct answer

Compute Engine

Your answer is incorrect

App Engine

Overall explanation

#1 is incorrect. Not all code is compatible with App Engine and some modifications might be necessary in order for it to work

#2 is incorrect. Not all code can be containerized and run on GKE and some modifications might be necessary in order for it to work

#3 is correct. Compute Engine offer the most flexibility and control, and works best in migration scenarios

Question 50 Correct

Which of the following services is NOT involved in the deployment process of Cloud Function?

Your answer is correct

GKE

Cloud Run

Cloud Storage

Overall explanation

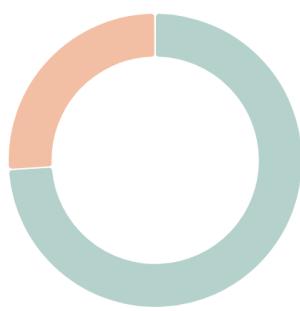
#1 is correct. GKE is not part of the deployment process of Cloud Function

#2 is incorrect. Code deployed to Cloud Function is stored in Cloud Storage

#3 is incorrect. Cloud Run is the runtime environment of Cloud Function

Sample Exam - Results

50 questions | 2 hours | 70% correct required to pass



Attempt 2: Passed! (70% required to pass)

74% correct (37/50)

1 hour 26 minutes

May 23, 2024 at 06:15 PM

[Review questions](#)

Correct Wrong Skipped/Unanswered